

# WEEK 11 LECTURE

## *CLASSIFICATION*

# OBJECTIVES

- Logistic Regression
- Poisson Regression
- Naive Bayes Classifier
- K-Nearest Neighbors
- Decision Tree
- Support Vector Machine

# LOGISTIC REGRESSION

## ➤ Pros:

- It is a very simple and efficient algorithm.
- The model has low variance.
- It provides probability score for observations.

## ➤ Cons:

- It's not good at handling a large number of categorical features.
- It assumes that the data is free of missing values and predictors are independent of each other.

# POISSON REGRESSION

- Poisson regression is often used for modeling count data. It is used to answer the questions such as what factors can predict the frequency of an event.

- The model for Poisson regression model is

$$\log y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p + \varepsilon$$

- The estimated regression equation is

$$\log \hat{y} = b_0 + b_1 x_1 + b_2 x_2 + \cdots + b_p x_p$$

- Requirement: The count (dependent) variable's variance is not greater than its mean.
- Coefficient interpretation: When  $x_i$  is increased by 1 (other variables are held constant),  $y$  is multiplied by  $e^{b_i}$ .

# POISSON REGRESSION (CONT.)

- How to determine the estimated coefficients  $b_i$ :

1) Consider the likelihood function

$$L(\beta_0, \beta_1, \dots, \beta_p) = \prod_{i=1}^n \frac{e^{-\lambda(x_i)} \lambda(x_i)^{y_i}}{y_i!}$$

where  $\lambda(x_i) = e^{\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}}$ .

2) Maximize the likelihood function  $L(\beta_0, \beta_1, \dots, \beta_p)$ .

# POISSON REGRESSION (CONT.)

## ➤ Pros:

- It does not predict negative values. It has a minimum value of 0. Thus, it is useful for distribution in which the mean or the most typical value is close to 0.
- The model is useful for skewed data with long tail. So, it is appropriate for counts of rare events.
- In general, for counts of events, it is a better model compared to the normal model.

## ➤ Cons:

- The primary problem is over-dispersion.

# NAIVE BAYES CLASSIFIER

- Naive Bayes algorithm is a classification technique based on the Bayes' theorem. It assumes the independence of independent variables.
- Naive Bayes classifier is popular, especially in text classification, spam filtering, and recommendation systems.
- Suppose that we wish to classify an observation into one of  $K$  classes,  $K \geq 2$ . Let  $\pi_k$  denotes the overall or prior probability that a randomly chosen observation comes from the  $k^{\text{th}}$  class. Let  $f_k(X) = P(X|Y = k)$  denote the density function of  $X$  for an observation that comes from the  $k^{\text{th}}$  class. Then Bayes' theorem states that

$$p_k(x) = P(Y = k|X = x) = \frac{\pi_k f_k(x)}{\sum_{i=1}^K \pi_i f_i(x)}$$

# NAIVE BAYES CLASSIFIER (CONT.)

- Assuming that within the  $k^{\text{th}}$  class, the  $p$  predictors are independent. This means that for  $k = 1, 2, \dots, K$ ,

$$f_k(x) = f_{k1}(x_1) \cdot f_{k2}(x_2) \cdot \dots \cdot f_{kp}(x_p)$$

where  $f_{kj}$  is the density function of the  $j^{\text{th}}$  predictor among observations in the  $k^{\text{th}}$  class. The posterior probability is then given by

$$p_k(x) = P(Y = k|X = x) = \frac{\pi_k \cdot f_{k1}(x_1) \cdot f_{k2}(x_2) \cdot \dots \cdot f_{kp}(x_p)}{\sum_{i=1}^K [\pi_i \cdot f_{i1}(x_1) \cdot f_{i2}(x_2) \cdot \dots \cdot f_{ip}(x_p)]}$$

for  $k = 1, 2, \dots, K$ .

- Estimate  $\pi_1, \dots, \pi_K$  as the proportion of training observations belonging to the  $k^{\text{th}}$  class, for  $k = 1, 2, \dots, K$ .



# NAIVE BAYES CLASSIFIER (CONT.)

- To estimate the density functions  $f_{kj}$  using training data  $x_{1j}, \dots, x_{nj}$ , consider a few options.
  - If  $X_j, j = 1, \dots, p$ , are quantitative, we can assume that  $X_j|Y = k \sim N(\mu_{jk}, \sigma_{jk}^2)$ , i.e. within each class, the  $j^{\text{th}}$  predictor is normally distributed.
  - If  $X_j, j = 1, \dots, p$ , are quantitative, we can also make a histogram for the observations of the  $j^{\text{th}}$  predictor within each class. Then we can estimate  $k_{kj}(x_j)$  as the fraction of the training observations in the  $k^{\text{th}}$  class that belong to the same histogram bin as  $x_j$ . Alternatively, we can use a kernel density estimator, which is essentially a smoothed version of a histogram.

## NAIVE BAYES CLASSIFIER (CONT.)

- If  $X_j, j = 1, \dots, p$ , are qualitative, then we can simply count the proportion of training observations for the  $j^{\text{th}}$  predictor corresponding to each class. For instance, suppose that  $X_j \in \{1, 2, 3\}$ , and we have 100 observations in the  $k^{\text{th}}$  class. Suppose that the  $j^{\text{th}}$  predictor takes on values of 1, 2, and 3 in 32, 55, and 13 of those observations, respectively. Then we can estimate  $f_{kj}$  as

$$\hat{f}_{kj}(x_j) = \begin{cases} 0.32 & \text{if } x_j = 1 \\ 0.55 & \text{if } x_j = 2 \\ 0.13 & \text{if } x_j = 3 \end{cases}$$

# NAIVE BAYES CLASSIFIER (CONT.)


## ➤ Pros:

- This algorithm is easy and works very fast.
- It can also be used to solve multi-class prediction problems as it's quite useful with them.
- This classifier performs better than other models with less training data if the assumption of independence of features holds.

## ➤ Cons:

- It assumes that all the features are independent. In real life, this is hard to hold true.
- It is also known as a bad estimator, so the probability outputs are not to be taken too seriously.

# NAIVE BAYES CLASSIFIER (CONT.)



**FIGURE 2.13.** A simulated data set consisting of 100 observations in each of two groups, indicated in blue and in orange. The purple dashed line represents the Bayes decision boundary. The orange background grid indicates the region in which a test observation will be assigned to the orange class, and the blue background grid indicates the region in which a test observation will be assigned to the blue class.


# K-NEAREST NEIGHBORS (KNN)

- KNN attempts to estimate the conditional distribution of  $Y$  given  $X$ , and then classify a given observation to the class with highest estimated probability. Given a positive integer  $K$  and a test observation  $x_0$ , the KNN classifier first identifies the  $K$  points in the training data that are closest to  $x_0$ , represented by the index set  $\mathcal{N}_0$ . It then estimates

$$P(Y = j|X = x_0) = \frac{1}{K} \sum_{i \in \mathcal{N}_0} I(y_i = j)$$


Finally, KNN classifies the test observation  $x_0$  to the class with the largest probability from the above probabilities.

# K-NEAREST NEIGHBORS (CONT.)



**FIGURE 2.14.** The KNN approach, using  $K = 3$ , is illustrated in a simple situation with six blue observations and six orange observations. Left: a test observation at which a predicted class label is desired is shown as a black cross. The three closest points to the test observation are identified, and it is predicted that the test observation belongs to the most commonly-occurring class, in this case blue. Right: The KNN decision boundary for this example is shown in black. The blue grid indicates the region in which a test observation will be assigned to the blue class, and the orange grid indicates the region in which it will be assigned to the orange class.

# K-NEAREST NEIGHBORS (CONT.)



**FIGURE 2.15.** The black curve indicates the KNN decision boundary on the data from Figure 2.13, using  $K = 10$ . The Bayes decision boundary is shown as a purple dashed line. The KNN and Bayes decision boundaries are very similar.

# K-NEAREST NEIGHBORS (CONT.)

## ➤ Pros:

- It is a simple and easy to implement.
- There is no need to make additional assumptions.
- It can be used for classification, regression and search.

## ➤ Cons:


- It gets slow as the number of data points and/or predictors/features increase.
- It assumes that the data is free of missing values and predictors are independent of each other.



# DECISION TREE

- Decision trees can be applied to both regression and classification problems.
- Roughly speaking, there are two steps in building a regression tree.
  - 1) Divide the predictor space — that is, the set of possible values for  $X_1, X_2, \dots, X_p$  — into  $J$  distinct and non-overlapping regions,  $R_1, R_2, \dots, R_J$ .
  - 2) For every observation that falls into the region  $R_j$ , we make the same prediction, which is simply the mean of the response values for the training observations in  $R_j$ .
- The construction of the regions  $R_1, R_2, \dots, R_J$  is beyond the scope of this course.

# DECISION TREE (CONT.)

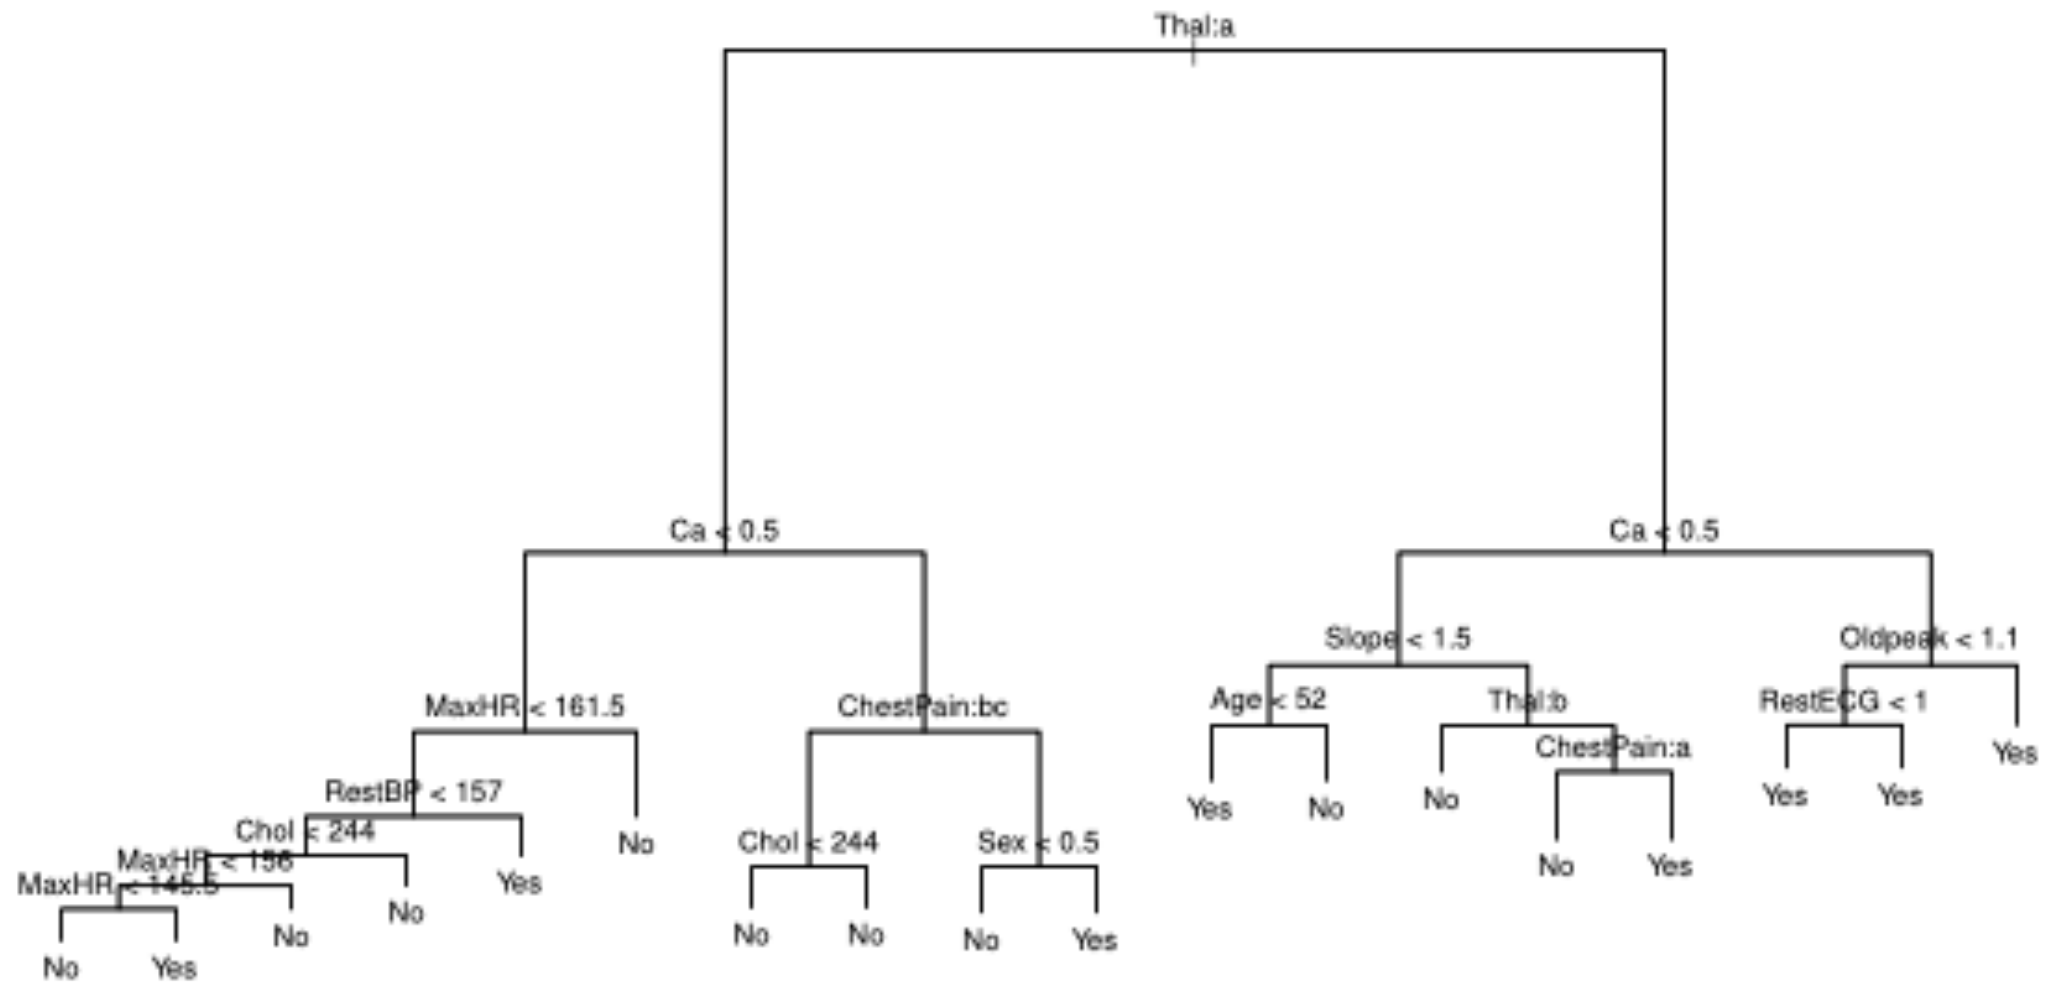


**FIGURE 8.4.** Regression tree analysis for the **Hitters** data. The unpruned tree that results from top-down greedy splitting on the training data is shown.

# DECISION TREE (CONT.)

- A classification tree is very similar to a regression tree, except that it is used to predict a qualitative response rather than a quantitative one. There are two main steps:
  - 1) Divide the predictor space — that is, the set of possible values for  $X_1, X_2, \dots, X_p$  — into  $J$  distinct and non-overlapping regions,  $R_1, R_2, \dots, R_J$ .
  - 2) For every observation that falls into the region  $R_j$ , we predict that each observation belongs to the most commonly occurring class of training observations in the region  $R_j$ .

# DECISION TREE (CONT.)

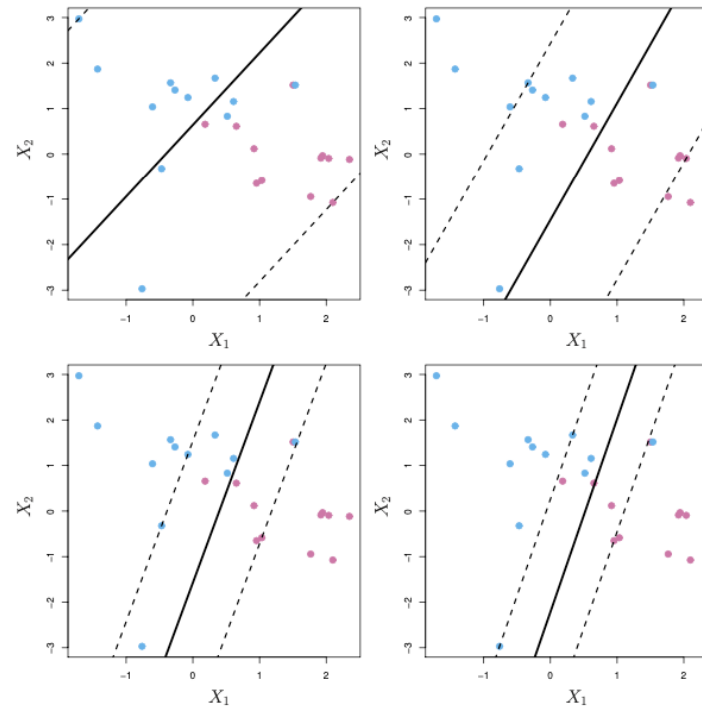


# SUPPORT VECTOR CLASSIFIER

- Support vector classifier finds a hyperplane in an  $N$ -dimensional space ( $N$  – the number of features) that distinctly classifies the data points. SVM can be used for both regression and classification tasks, but it is widely used in classification.
- The hyperplane is chosen to correctly separate most of the training observations into the two classes, but it may misclassify a few observations. It is the solution to the optimization problem

$$\begin{aligned} & \underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n, M}{\text{maximize}} && M \\ & \text{subject to} && \sum_{j=1}^p \beta_j^2 = 1, \\ & && y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i), \\ & && \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C, \end{aligned}$$


# SUPPORT VECTOR CLASSIFIER (CONT.)



**FIGURE 9.7.** A support vector classifier was fit using four different values of the tuning parameter  $C$  in (9.12)–(9.15). The largest value of  $C$  was used in the top left panel, and smaller values were used in the top right, bottom left, and bottom right panels. When  $C$  is large, then there is a high tolerance for observations being on the wrong side of the margin, and so the margin will be large. As  $C$  decreases, the tolerance for observations being on the wrong side of the margin decreases, and the margin narrows.


# SUPPORT VECTOR CLASSIFIER (CONT.)

- The support vector classifier is a natural approach for classification in the two-class setting, if the boundary between the two classes is linear. However, in practice we are sometimes faced with non-linear class boundaries.



**FIGURE 9.8.** Left: The observations fall into two classes, with a non-linear boundary between them. Right: The support vector classifier seeks a linear boundary, and consequently performs very poorly.

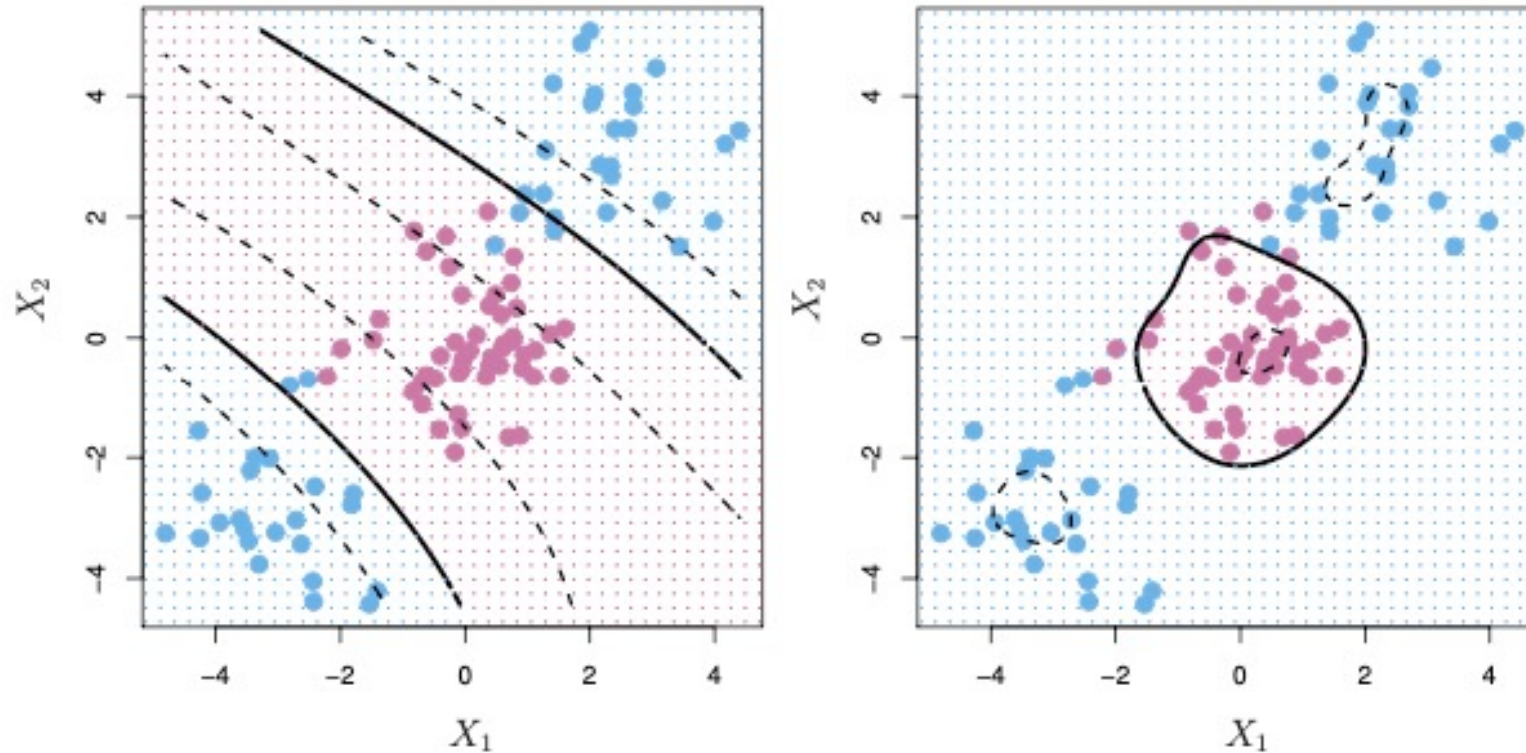
# SUPPORT VECTOR MACHINE (SVM)



Source: <https://www.analyticsvidhya.com/blog/2021/10/support-vector-machinessvm-a-complete-guide-for-beginners/>



# SUPPORT VECTOR MACHINE (CONT.)



**FIGURE 9.9.** Left: An SVM with a polynomial kernel of degree 3 is applied to the non-linear data from Figure 9.8, resulting in a far more appropriate decision rule. Right: An SVM with a radial kernel is applied. In this example, either kernel is capable of capturing the decision boundary.

# SUPPORT VECTOR MACHINE (CONT.)

## ➤ Pros:

- It work well when there is a clear margin of separation between classes.
- It is more effective in high-dimensional spaces.

## ➤ Cons:

- It is not good for large data sets.
- It is not good when classes are overlapping.